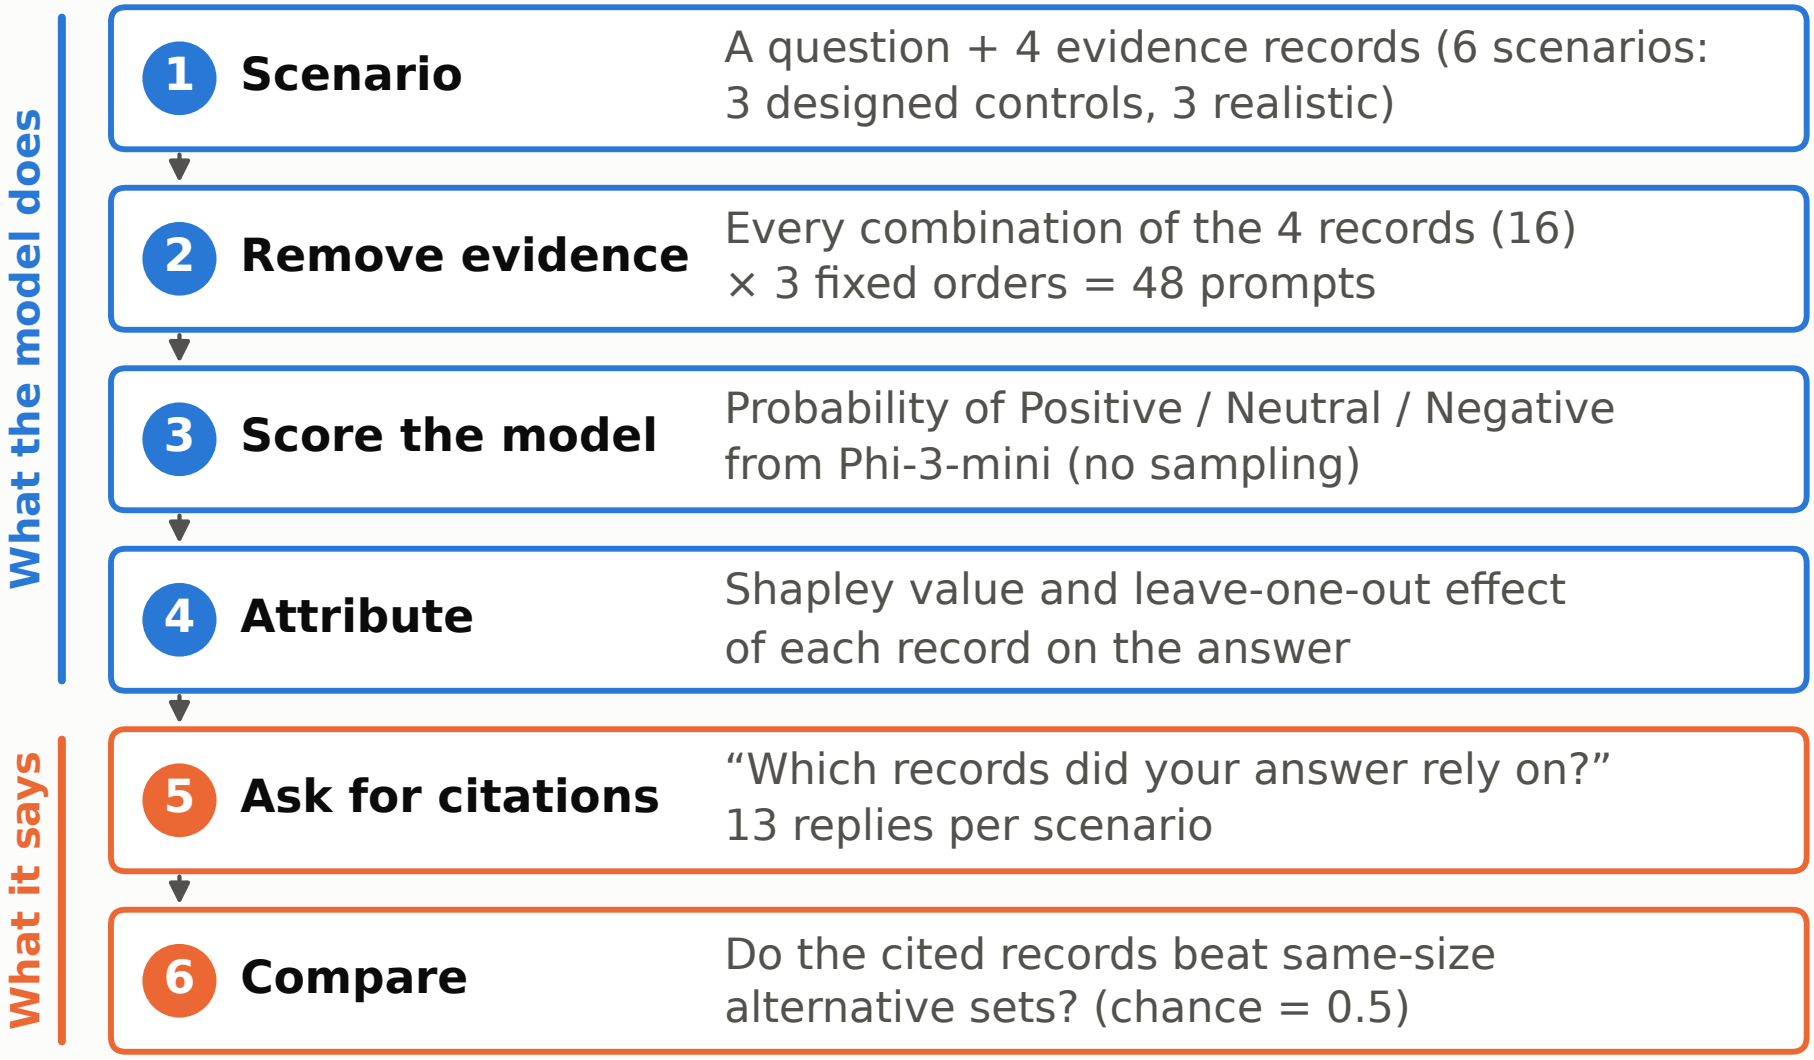
What the model does
What it says
1 Scenario
A question + 4 evidence records (6 scenarios: 3 designed controls, 3 realistic)
2 Remove evidence
Every combination of the 4 records (16) × 3 fixed orders = 48 prompts
3 Score the model
Probability of Positive / Neutral / Negative from Phi-3-mini (no sampling)
4 Attribute
Shapley value and leave-one-out effect of each record on the answer
5 Ask for citations
“Which records did your answer rely on?” 13 replies per scenario
6 Compare
Do the cited records beat same-size alternative sets? (chance = 0.5)
Validation infrastructure checks · designed controls · independent re-analysis
Post hoc content-only citer from the design or rubric; never sees the model